

```
In [3]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
```

## Pre processing img data

```
In [4]: train_dir = "cifar-10-img/cifar-10-img/train"
test_dir = "cifar-10-img/cifar-10-img/test"

train_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
)

test_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
)

# here batch_size is the number of images in each batch
train_batch_size = 5000
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(32, 32),
    batch_size=train_batch_size,
    class_mode='categorical'
)
test_batch_size = 1000
test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(32, 32),
    batch_size=test_batch_size,
    class_mode='categorical'
)
```

Found 40079 images belonging to 10 classes.  
Found 9921 images belonging to 10 classes.

## Selecting only first batch with 5000 images as train and test data

```
In [5]: x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]

print(len(x_train))
print(len(x_test))
```

5000  
1000

### a. Load in a pre-trained CNN model trained on a large dataset

```
In [6]: # Load VGG16 without top layers
weights_path = "vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5"
base_model = VGG16(weights=weights_path, include_top=False, input_shape=(32, 32, 3))
```

## b. Freeze parameters (weights) in model's lower convolutional layers

```
In [7]: for layer in base_model.layers:
    layer.trainable = False
```

## c. Add custom classifier with several layers of trainable parameters to model

```
In [8]: x = Flatten()(base_model.output)
x = Dense(256, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
x = Dense(256, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)
# Compile the model
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy'])
```

## d. Train classifier layers on training data available for task

```
In [9]: # Train the model
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_te
```

```

Epoch 1/10
79/79 65s 767ms/step - accuracy: 0.2858 - loss: 1.9748 - val_accuracy: 0.3860 - val_loss: 1.6505
Epoch 2/10
79/79 60s 761ms/step - accuracy: 0.4368 - loss: 1.5820 - val_accuracy: 0.4820 - val_loss: 1.4761
Epoch 3/10
79/79 59s 746ms/step - accuracy: 0.4848 - loss: 1.4490 - val_accuracy: 0.5050 - val_loss: 1.4197
Epoch 4/10
79/79 59s 742ms/step - accuracy: 0.5258 - loss: 1.3451 - val_accuracy: 0.5160 - val_loss: 1.3857
Epoch 5/10
79/79 60s 761ms/step - accuracy: 0.5300 - loss: 1.3036 - val_accuracy: 0.5260 - val_loss: 1.3878
Epoch 6/10
79/79 60s 764ms/step - accuracy: 0.5596 - loss: 1.2402 - val_accuracy: 0.5340 - val_loss: 1.3614
Epoch 7/10
79/79 81s 751ms/step - accuracy: 0.5696 - loss: 1.1968 - val_accuracy: 0.5390 - val_loss: 1.3443
Epoch 8/10
79/79 60s 756ms/step - accuracy: 0.5964 - loss: 1.1397 - val_accuracy: 0.5430 - val_loss: 1.3512
Epoch 9/10
79/79 59s 750ms/step - accuracy: 0.6054 - loss: 1.1018 - val_accuracy: 0.5360 - val_loss: 1.3331
Epoch 10/10
79/79 59s 743ms/step - accuracy: 0.6202 - loss: 1.0582 - val_accuracy: 0.5320 - val_loss: 1.3575

```

Out[9]: <keras.src.callbacks.history.History at 0x21ddcd1160>

### e. Fine-tune hyper parameters and unfreeze more layers as needed

```

In [10]: base_model = VGG16(weights=weights_path, include_top=False, input_shape=(32, 32, 3))
# freeze all layers first
for layer in base_model.layers:
    layer.trainable = False
# unfreeze last 4 layers of base model
for layer in base_model.layers[len(base_model.layers) - 4:]:
    layer.trainable = True
# fine-tuning hyper parameters
x = Flatten()(base_model.output)
x = Dense(256, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
x = Dense(512, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)
# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy',

```

```
# training fine tuned model
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_te

Epoch 1/10
79/79 ━━━━━━━━━━ 105s 1s/step - accuracy: 0.1576 - loss: 2.2179 - val_accur
accuracy: 0.1740 - val_loss: 2.0656
Epoch 2/10
79/79 ━━━━━━━━━━ 138s 1s/step - accuracy: 0.2368 - loss: 1.9183 - val_accur
accuracy: 0.3770 - val_loss: 1.5908
Epoch 3/10
79/79 ━━━━━━━━━━ 153s 1s/step - accuracy: 0.4748 - loss: 1.4207 - val_accur
accuracy: 0.4490 - val_loss: 1.6163
Epoch 4/10
79/79 ━━━━━━━━━━ 94s 1s/step - accuracy: 0.5984 - loss: 1.1576 - val_accur
accuracy: 0.5770 - val_loss: 1.2232
Epoch 5/10
79/79 ━━━━━━━━━━ 143s 1s/step - accuracy: 0.6732 - loss: 0.9533 - val_accur
accuracy: 0.5870 - val_loss: 1.2751
Epoch 6/10
79/79 ━━━━━━━━━━ 140s 1s/step - accuracy: 0.7122 - loss: 0.8254 - val_accur
accuracy: 0.6100 - val_loss: 1.2426
Epoch 7/10
79/79 ━━━━━━━━━━ 148s 1s/step - accuracy: 0.7548 - loss: 0.7015 - val_accur
accuracy: 0.5850 - val_loss: 1.3624
Epoch 8/10
79/79 ━━━━━━━━━━ 138s 1s/step - accuracy: 0.7888 - loss: 0.6338 - val_accur
accuracy: 0.6230 - val_loss: 1.3001
Epoch 9/10
79/79 ━━━━━━━━━━ 84s 1s/step - accuracy: 0.8150 - loss: 0.5188 - val_accur
accuracy: 0.6020 - val_loss: 1.3733
Epoch 10/10
79/79 ━━━━━━━━━━ 83s 1s/step - accuracy: 0.8332 - loss: 0.5054 - val_accur
accuracy: 0.5870 - val_loss: 1.8079
```

Out[10]: <keras.src.callbacks.history.History at 0x21de57b2350>


In [11]: `import matplotlib.pyplot as plt  
predicted_value = model.predict(x_test)`

32/32 ━━━━━━━━ 10s 302ms/step

In [12]: `labels = list(test_generator.class_indices.keys())`

In [13]: `n = 890  
plt.imshow(x_test[n])  
print("Predicted: ", labels[np.argmax(predicted_value[n])])  
print("Actual: ", labels[np.argmax(y_test[n])])`

Predicted: truck  
Actual: truck



In [ ]: